

# CRM Sync — What Traditional CRMs Miss

For business leaders, investors, and operations teams evaluating CRM Sync against Salesforce, HubSpot, and Klaviyo. **Traditional CRMs were built for a pre-AI, pre-privacy world** — centralized databases, batch syncs, vendor-locked identity, GDPR bolted on afterward. CRM Sync is consent-first and AI-native from the ground up.

For: Business leaders, investors, and operations teams evaluating CRM Sync against Salesforce, HubSpot, and Klaviyo

Date: 2026-05-19 · Updated: 2026-07-19

---

## Summary

---

Traditional CRMs were built for a pre-AI, pre-privacy world: centralized databases, batch syncs, vendor-locked identity, and GDPR bolted on as an afterthought. CRM Sync is designed from the ground up for real-time, consent-first, AI-native customer data orchestration across composable services.

The endgame is visible in two steps. The **GA4 pivot** moves events and identity into infrastructure you own (server-side Measurement Protocol, the free BigQuery export, your own project and IAM). The **UCP turn** makes commerce agent-addressable. At that threshold, organizations pivoting to GA4 discover the CRM was never the requirement — **consent-governed identity was**, and by then they already own it.

---

# 1. Consent-First Architecture

---

Salesforce and HubSpot store consent as a checkbox field. CRM Sync treats consent as an infrastructure gate — no data moves to any channel until consent is verified, with a full audit trail (timestamp, method, browser, session ID).

Traditional CRMs rely on the marketing team to check the box. CRM Sync enforces it at the infrastructure level.

**The banner most enterprises run is still tuned for the page-view era.** UA-generation CMP deployments — the OneTrust banner wired up years ago — record a page-view consent and stop there, and household-name enterprises are still running them because nobody owns the update — these are the same organizations whose UA→GA4 migrations became multi-year programs that broke reporting on the way. An org that can't move its own analytics tag without disintegrating is not going to move consent on its own. The core problem isn't the banner: **they have no place to put or log real-time data** — no ledger for an event to land in — and the consulting teams keep selling them UA-shaped, page-view software, so the gap renews itself with every engagement. This is where the AI component actually works: it **folds into the tools already running** — nothing ripped out, the ledger installed beside them — and then upgrades the organization to more efficient, easier-to-use data management from there. But AI predictive audiences explicitly require **Consent Mode v2**: no v2 signals, no modeling, no predictive audiences — and no defensible record. Most orgs treat CMv2 as optional. Then they're surprised when they get sued. The ones that treat it as infrastructure never have that day.

## 2. Real-Time Multi-Channel Sync (Not Batch ETL)

---

Klaviyo syncs customer segments on a schedule (hourly/daily). HubSpot workflows trigger sequentially. CRM Sync writes to all 7 services in the same request — database, website CMS, online store, analytics, customer data platform, email, and consent log — with zero delay.

No "sync lag" where a customer opts out of marketing but still receives an email because the batch hasn't run yet.

## 3. No Vendor Lock-In on Identity

---

Salesforce owns your customer identity inside their platform. Exporting means CSV dumps and field mapping. CRM Sync stores identity in your own database (Xano) with an open schema — you control the data, not the CRM vendor.

## 4. Edge-Native, Not Cloud-Monolith

---

HubSpot and Salesforce run from centralized data centers. CRM Sync runs on Cloudflare's edge network (300+ cities), so auth checks and consent verification happen in milliseconds, close to the user. No cold starts, no server scaling decisions.

## 5. Per-Tenant Isolation by Default

---

Salesforce multi-org requires Enterprise licensing. HubSpot business units are a premium add-on. CRM Sync isolates every store at the data key level — separate credentials, separate admin keys, separate config — on the base plan.

## 6. AI Agent Security Boundaries

---

Traditional CRMs give API keys broad access. A Salesforce Connected App with full API scope can read and write everything. CRM Sync scopes AI tools to specific credentials, logs every agent action with source attribution, and masks all secrets in diagnostic views.

No traditional CRM has a concept of "this API key is for an AI agent and should be short-lived and auditable."

## 7. Composable, Not Monolithic

---

HubSpot bundles CMS, CRM, email, and analytics into one product — you pay for all of it. CRM Sync is a thin orchestration layer that connects best-of-breed tools you already use (Shopify, Webflow, GA4, Adobe, Resend). Swap any service without re-platforming.

## 8. GDPR as Infrastructure, Not Add-On

---

Salesforce charges extra for Privacy Center. HubSpot GDPR tools are limited to consent banners and deletion requests. CRM Sync has GDPR baked into the data model: Article 7 consent provenance, Article 15 data export, Article 17 right-to-erasure, Article 20 portability — all as standard endpoints with cryptographically verified compliance webhooks.

## 9. Server-Side Analytics Without Tag Managers

---

Klaviyo and HubSpot depend on client-side JavaScript for tracking, which ad blockers defeat. CRM Sync pushes user properties and events to GA4 via server-side Measurement Protocol — immune to ad blockers, consistent data regardless of browser.

## 10. Transparent Pricing Without Per-Contact Scaling

---

HubSpot charges by contact tier (\$800/mo at 2,000 contacts on Professional). Klaviyo charges per profile. Salesforce charges per user seat plus per-feature add-ons. CRM Sync charges flat, published prices — no per-contact, per-profile, or per-seat penalty for growing your customer base.

**What you're buying is a build, not a tenancy.** CRM Sync is a build with AI integration tooling: run it alongside the standard CRM you already have, or mirror that CRM's data through the sync layer and migrate off it entirely — the mirror becomes the system you own. Either path uses the same connectors and the same consent substrate; there is no fork-lift moment where the vendor decides your schema.

**Your data is yours. AI helps you keep it secure.** We generate a Worker — AI customizes it for you — and you use it where you need it. The deliverable is an edge Worker on Cloudflare Workers: AI-secure by construction (scoped credentials, audited agent actions, masked secrets), with the Verified Trust network — published keys, signed certificates, offline-verifiable records — reinforcing the system you already have rather than replacing it. Built to enterprise standards with flexibility in mind: it meets your stack where it stands.

This is also why AI-paired system development stays **front-end agnostic** here: none of the security lives in the page. Every read and write passes through surgical, row-based AI functions at the edge — each call scoped to the row it touches and the mandate it carries, consent checked at that row, the result ledgered. Swap the front end and nothing security-relevant moves.

**The Worker that secures the configurator is shared infrastructure — included.** The \$90 configurator runs against the shared edge Worker: the function/security brain — consent gates, entitlements, ledgers, the Verified Trust network — operated for every tenant at once. Want your own brain? A **dedicated Worker** on your own zone, with your own keys, your own ledgers, your own trust network, generated and customized for you — that is the Enterprise engagement, priced by scope with the private-customization menu below. The shape is familiar: WordPress on shared hosting versus your own private VM — identical software, different isolation, and you graduate when the stakes say so.

**The key strategy works because there is no "store your data somewhere else" fee.** Identity and records live in your own Shopify, Xano, and BigQuery under your own keys — per-tenant isolation is key-level, not a licensing tier. The price buys the build, never a place to keep what was already yours.

**Single-fee services** (one-time, fixed price):

| SERVICE                                                                                                  | PRICE   |
|----------------------------------------------------------------------------------------------------------|---------|
| <u>Download App</u> — the configurator: use as you wish, on as many Webflow/Shopify instances as you run | \$90    |
| <u>Verified Transaction Layer — License</u>                                                              | \$90    |
| <u>Firmware SBOM — Per-Image Analysis</u>                                                                | \$250   |
| <u>Firmware Security — ceremony-signed distribution</u>                                                  | \$1,900 |
| <u>CRA Readiness Assessment</u> — fixed-fee working session, all ten CRA documents in order              | \$900   |

**CRA compliance line** (what an EU Cyber Resilience Act posture costs here, versus a platform's enterprise-compliance add-on tier):

| CRA PRODUCT                                     | PRICE            |
|-------------------------------------------------|------------------|
| <u>SBOM Registry — App Coverage</u>             | \$49/mo          |
| <u>Firmware SBOM — Per-Image Analysis</u>       | \$250 one-time   |
| <u>Firmware Security — vault + certificates</u> | \$1,900 one-time |
| <u>CRA Readiness Assessment</u>                 | \$900 one-time   |

Distribution and creator rails (Channel Publish, 3D Publisher) are \$49 each. Every price above is the live store price — the catalog is the price list.

**Private customization** (scoped engagements on the same substrate — priced by scope, not by seat):

- **ERP integration** — your ERP's records join the universal revenue ledger on the session key, so order, tax, consent, and ERP references reconcile in one place instead of four exports. This is the gap acquisitions keep rediscovering: an ERP that doesn't resolve to the transaction record.
- **RMA / returns lifecycle** — returns as first-class records with verifiable state, joined to the same ledger; the "customer service system that doesn't link to the warehouse" failure mode, closed.
- **Fraud linkage** — fraud signals bound to the CRM identity spine rather than a sidecar tool, so a flagged actor is flagged everywhere consent allows.
- **Cross-border / market rails** — per-market pricing, tax, and payment rails (card, regional wallets) under one market-of-sale record per transaction, with EU price-reference (Omnibus) discipline built in.
- **Add-ons** — anything the composable layer reaches: warehouse/WMS hooks, channel publishing, custom consent surfaces, agent-payment mandates.

Because the substrate is the same consent-gated event bus, a custom add-on inherits the audit trail instead of escaping it.

- **Bring your own interface** — the data plane doesn't care what renders it. Managed CMS: Webflow, Em Dash, WordPress, Drupal, AEM, Salesforce Experience Cloud. File-system frameworks: Astro, Next, Nuxt, Svelte. The connectors and embeds are interface-agnostic, so the presentation layer is your choice — and replaceable without touching the data.

Standard CRMs sell these as enterprise editions of themselves. Here they're builds on infrastructure you already own by the time you need them.

## 11. Attio and the AI-Native CRMs: the Hotel California Problem

---

Newer, AI-native CRMs (Attio is the one evaluators search most) fix the UI and the data model — and change nothing about the shape that matters. The record still accumulates inside the vendor's workspace, per-seat billing still scales with the team, and consent is still a field you maintain rather than a timestamped object the infrastructure enforces. Even teams that go "headless CRM" hit the same wall: the painful part is realizing later that **you can never leave** — every workflow points at the vendor's IDs and the exit is a CSV plus a re-keying project. The bridge is going to be needed either way: an AI tool that mirrors the vendor's data into systems you own while everything keeps running. Install it on day one, not on the day you want out. Store page for evaluators: [Attio Alternative](#).

## 12. The Migration Rail: CRM → Shopify → GA4 / UCP

---

CRMs weren't built to include consent in real time. You will need to **update or migrate** — CRM Sync does both. The same connectors that let you coexist with a CRM are a **data-harvesting rail for leaving it** — three stages, no big-bang:

1. **CRM → Shopify.** Point the sync layer at the CRM in mirror mode. Profiles, consent-relevant records, and transaction history harvest into Shopify customer objects and your own identity spine — the CRM keeps running while its monopoly on the record ends.
2. **Shopify → GA4.** Events and identity flow server-side (Measurement Protocol, Consent Mode v2) into GA4 and the free BigQuery export — infrastructure you own, keyed by consent-gated `user_id`.
3. **GA4 → UCP.** The same plane publishes agent-eligible offers over the Universal Commerce Protocol. Humans arrive through Google; agents arrive through UCP; one consent substrate governs both.

**Logging on change is the difference.** Every change — price, consent, campaign, discount, A-B test — lands in a ledger the business stakeholder owns: their own business notebook, exportable as a CSV, visible only to them. That does two things at once: it adds **security** (a private, tamper-evident record nobody else can rewrite) and it adds **utility to AI-managed data** (the AI works against a ledger the human can always read, check, and keep).

**The benefit of direct-to-user data with consent is real time.** A CRM's copy of the customer is always a replica — captured somewhere else, synced later, consent checked after the fact. Direct-to-user capture inverts that: the record is written at the moment of the event, from the user themselves, with consent evaluated in the same request. There is no sync lag to apologize for and no

replica to reconcile — the first copy is already yours, already consented, already in the systems you keep.

## 13. Warehouse Identity: What a BigQuery ID Costs, and What Consent Unlocks

---

Salesforce sells warehouse-grade identity as Data Cloud — consumption credits on top of platform licensing. HubSpot gates warehouse sync behind Operations Hub. Klaviyo offers CSV exports. CRM Sync takes the commodity path: a consent-gated `user_id` flows through GA4's free BigQuery export into a Google Cloud project you own.

**The cost side is commodity infrastructure.** The GA4 daily export to BigQuery is free. BigQuery's free tier (10 GiB storage, 1 TiB of queries per month) covers a typical store end-to-end; past that you pay published on-demand rates (about \$6.25 per TiB queried, about \$0.02 per GiB-month stored) — pennies per month at DTC volumes, in your own project under your own IAM. There is no per-contact or per-profile multiplier on identity.

**The benefit side only exists with consent.** The `user_id` is attached only when the consent gate has passed (Consent Mode v2). A non-consented session still produces aggregate analytics, but never an identity-bearing row. Everything downstream inherits that property: retargeting audiences built from newsletter signups, revenue-weighted Smart Bidding signals pushed through Data Manager, predicted-LTV models trained in BQML. Each is compliant by construction, because a non-consented ID was never emitted in the first place.

The asymmetry is the point: carrying the ID costs almost nothing, while the consented benefit is a first-party audience and bidding asset that survives ad blockers, cookie deprecation, and ad-platform policy churn. Traditional CRMs

invert this — you pay platform rates for identity whether or not you can lawfully activate it.

---

# Comparison Table

| CAPABILITY                                                  | SALESFORCE                  | HUBSPOT               | KLAVIYO                | CRM SYNC                                                                                                              |
|-------------------------------------------------------------|-----------------------------|-----------------------|------------------------|-----------------------------------------------------------------------------------------------------------------------|
| Consent enforcement                                         | Checkbox field              | Banner + field        | Opt-in list            | Infrastructure gate with audit trail                                                                                  |
| Sync latency                                                | Batch / workflow            | Sequential workflow   | Hourly/daily           | Real-time (same request)                                                                                              |
| Identity ownership                                          | Vendor-locked               | Vendor-locked         | Vendor-locked          | Your database (Xano)                                                                                                  |
| Infrastructure                                              | Centralized cloud           | Centralized cloud     | Centralized cloud      | Edge network (300+ cities)                                                                                            |
| Multi-tenant isolation                                      | Enterprise add-on           | Premium add-on        | Not available          | Included on base plan                                                                                                 |
| AI agent security                                           | Broad API scope             | Broad API scope       | Broad API scope        | Scoped credentials + audit trail                                                                                      |
| Architecture                                                | Monolithic platform         | Monolithic platform   | Email-focused platform | Composable orchestration                                                                                              |
| GDPR compliance                                             | Privacy Center add-on       | Basic tools           | Limited                | Built into data model                                                                                                 |
| Analytics tracking                                          | Client-side JS              | Client-side JS        | Client-side JS         | Server-side (ad-block proof)                                                                                          |
| Warehouse identity                                          | Data Cloud add-on (credits) | Operations Hub add-on | CSV export             | Consent-gated <code>user_id</code> in your BigQuery (free GA4 export)                                                 |
| Evidence ledger (Omnibus pricing · firmware/SBOM · consent) | No                          | No                    | No                     | <b>One ledger — unique to CRM Sync:</b> price evidence, shipped-software records, and Consent Mode v2 logged together |

| CAPABILITY               | SALESFORCE         | HUBSPOT          | KLAVIYO                       | CRM SYNC                                                                                                         |
|--------------------------|--------------------|------------------|-------------------------------|------------------------------------------------------------------------------------------------------------------|
| Geo-dimensioned evidence | No                 | No               | Profile location filters only | <b>Every event carries ISO country, currency, language, location</b> — per-market evidence, not a segment filter |
| Pricing model            | Per seat + add-ons | Per contact tier | Per profile                   | Flat published prices — a build you own (\$90 app; fixed-fee services)                                           |