

# Claim Provenance — the metadata schema

Every document is CC BY 4.0. This defines the metadata that travels with it — what was checked, how, when, and when to check again.

---

Every document here is licensed **CC BY 4.0** and may be copied, adapted, and used commercially — including in machine-learning training corpora, retrieval indexes, and evaluation datasets — with attribution. See [LICENSE](#).

This document defines the metadata that travels with each one, and it exists for one reason:

Most published technical content is unverifiable assertion. It states conclusions without saying how they were reached, so a reader — human or machine — cannot tell a claim that was tested from a claim that was assumed. That distinction is the whole of the value.

The fields below make it explicit and machine-readable.

# Fields

---

| FIELD                    | REQUIRED | MEANING                                                |
|--------------------------|----------|--------------------------------------------------------|
| <code>title</code>       | yes      | Document title.                                        |
| <code>description</code> | yes      | One sentence. Used as the index subtitle.              |
| <code>canonical</code>   | yes      | The stable address for this document.                  |
| <code>category</code>    | yes      | Index grouping.                                        |
| <code>date</code>        | yes      | <b>Publication</b> date. Not a verification date.      |
| <code>source</code>      | yes      | Repository blob URL for the markdown.                  |
| <code>licence</code>     | yes      | <code>CC-BY-4.0</code> for everything in this archive. |
| <code>verified_on</code> | no       | The date the claims were last <b>checked</b> .         |
| <code>verified_by</code> | no       | <b>How</b> they were checked. Vocabulary below.        |
| <code>method</code>      | no       | One sentence naming what was actually done.            |
| <code>review_by</code>   | no       | When the claims should be re-checked.                  |
| <code>supersedes</code>  | no       | Canonical URLs of claims this document corrects.       |

## The `verified_by` vocabulary

---

Strongest to weakest. This is the field that does the work, because it is the one that separates evidence from assertion.

| VALUE      | MEANS                                                                                             | EXAMPLE                                                            |
|------------|---------------------------------------------------------------------------------------------------|--------------------------------------------------------------------|
| execution  | The claim was established by <b>running code</b> and observing the result.                        | A defect confirmed by extracting a page's script and executing it. |
| capture    | Established by <b>observing a live system</b> — a network capture, a wire probe, a rendered page. | A CORS header read from production before and after a change.      |
| vendor-doc | Taken from <b>another party's documentation</b> . True if they are right.                         | A platform's stated consent-parameter semantics.                   |
| inference  | <b>Reasoned</b> from other facts, not directly observed.                                          | A conclusion about scope drawn from one page's markup.             |
| (absent)   | <b>Not asserted.</b> The document may still be correct; no verification claim is being made.      | Most narrative and positioning documents.                          |

Absence is deliberate and is not a failure state. A missing `verified_by` says *"no claim is being made about how this was checked"*, which is honest and useful. Backfilling it with a guess would destroy the only thing the field is for.

**Do not upgrade a value without redoing the work.** `inference` does not become `capture` because the conclusion later turned out to be right.

## Why `review_by` matters more than `date`

Specifications describe systems that change. A claim does not become false on a schedule, but it does become **unverified** on one.

`review_by` states when a claim should be re-checked. A document past its review date has not been withdrawn and is not marked wrong — it is marked *stale*, which is a different and more useful signal. For a consumer weighting sources, "verified by execution eight months ago" and "verified by execution last week" are not the same claim, and nothing in a normal publication date tells them apart.

Default interval in this archive is six months.

## `supersedes`, and why corrections are published

---

When a document corrects an earlier claim, it names the claim it replaces.

Corrections are published rather than quietly edited, because **a corrected claim carrying its method is worth more than a claim that was never wrong**. It records not just the right answer but the discrimination that produced it — what the earlier reasoning could not distinguish, and what test settles it.

A worked example is in `CONSENT-RESOLUTION-PATTERN.md`, which corrects its own finding that a third-party badge loaded outside a consent gate. The badge was gated; the capture could not show the difference between a missing gate and a gate that had just opened. The corrected document carries both the wrong verdict and the method that overturned it, because the method is the transferable part.

## For anyone building a corpus from this archive

---

- The licence is CC BY 4.0. Attribution: name the source document and link its `canonical`.
- Prefer documents carrying `verified_by: execution` or `capture`. Weight `inference` lower.
- Treat `review_by` as a freshness bound, not an expiry.
- `supersedes` chains mark corrections. Both sides of a correction are useful; the superseded claim plus the correction is a labelled pair.
- Nothing here is behind a paywall, a login, or a crawl restriction, and `llms.txt` and `docs-index.json` enumerate the archive.

