

Process Management Guide — Webflow · Xano · Cloudflare · Shopify

When you operate across independent platforms, the most dangerous move is letting one of them be both your **source of truth** and your **runtime dependency**. The June 3 2026 global Shopify outage — and Klaviyo's failure the next day — are the case studies. This defines the role each platform plays, how data flows between them, and the runbook that keeps you running when a platform you don't control goes dark.

Audience: Operations, platform engineering, and business stakeholders running a multi-system commerce stack

Scope: Public guide. No secrets, credentials, or source code. Process, architecture, and automation recommendations only.

Date: 2026-06-08

0. TL;DR

When you operate across **multiple independent platforms**, the most dangerous decision you can make is to let any one of them be both your **source of truth** and your **runtime dependency**. The June 3, 2026 global Shopify outage — where storefronts worldwide resolved to *"This store does not exist"* — is the case study that proves it. And the very next day, **June 4, Klaviyo went down too** (logins, APIs, event ingestion, and campaign/flow sends), proving the broader point: *this is not a Shopify problem — every platform you don't control will eventually have a bad day, sometimes back-to-back*. This guide defines the roles each platform should play, how data should move between them (the

"ORM / Content / Data Layer" split), and the operating process that keeps you running when a platform you don't control goes dark.

One-line recommendation: *Shopify is a channel, not your database. Xano is the source of truth, Cloudflare is the resilient runtime, Webflow is the presentation layer, and Shopify is one of several destinations that data flows **to** — never the thing everything else flows **from**.*

1. The four platforms and the role each should play

PLATFORM	CORRECT ROLE	WHAT IT MUST NOT BECOME
Shopify	Commerce channel — catalog, checkout, payments, POS. A <i>destination</i> and a transaction engine.	The single source of truth for your catalog, customers, or content. A hard runtime dependency for pages that must stay up.
Xano	Data layer / source of truth — the relational system of record for catalog, customers, consent, orders, and sync state.	A passive cache. If Xano only mirrors Shopify, you have no independent truth to recover from.
Cloudflare	Resilient runtime — Worker that routes requests, verifies identity, enforces consent, caches reads at the edge, and absorbs upstream outages.	A thin proxy that fails the moment Shopify fails.
Webflow	Content / presentation layer — marketing pages, CMS content, and the Designer Extension your team operates.	A second uncontrolled source of truth that drifts from Xano.

The recurring failure mode across every team that has been burned: they let **Shopify be the database**. Then the day Shopify is unreachable, *everything* — storefront, content, internal tooling, analytics — goes down with it. The architecture below is designed so that a total Shopify outage degrades you to "checkout temporarily unavailable" instead of "business does not exist."

2. Case study – the June 3, 2026 global Shopify outage

2.1 What happened

On the morning of **June 3, 2026**, Shopify suffered a **global outage** affecting merchants and customers worldwide. Per Shopify's public status updates (reported by PYMNTS):

TIME (EDT)	STATUS
9:27 a.m.	Shopify announced issues across one or more functions: admin, checkout, storefronts, Retail POS, and support access
10:37 a.m.	Identified the problem; reported recovery from mitigation efforts
11:31 a.m.	Declared resolved
3:13 p.m.	Final note: <i>"This issue has been resolved, and we are continuing to monitor"</i>

The **officially reported window was roughly two hours**, with monitoring continuing into the afternoon. Many merchants reported **degraded and intermittent impact lasting much of the day** – so plan for the worst case (a multi-hour to all-day disruption), not the press-release best case. DownDetector logged **3,000+ reports**, spiking before 9 a.m. EDT.

2.2 The error that made it worse

During the outage, visitors to affected stores saw:

"This store does not exist" – accompanied by a Shopify advertisement.

This is the part that turns a vendor incident into *your* reputational incident. The error doesn't say "temporarily down." It tells your customers your business **no longer exists**, and advertises the platform on your storefront's grave. Merchants publicly demanded Shopify swap it for an honest outage notice:

"Change the 'This store does not exist' page to a page that says 'We are experiencing an outage — we'll be back soon.' Having our stores resolve to that page is insane."

2.3 The lessons (these drive every recommendation below)

1. **You do not control your platform's uptime, and you do not control its error page.** When Shopify is down, customers see Shopify's chosen message about *your* brand — and it can be actively damaging.
2. **A platform outage is a single point of failure if you let it be one.** If your homepage, your content, and your internal tools all depend on a live Shopify response, one outage takes out everything at once.
3. **"Resolved" for the vendor ≠ "recovered" for you.** Caches, webhooks, and queued events need to drain and reconcile after the platform returns. Budget for a recovery tail.
4. **Read availability and write availability are different problems.** You can keep *browsing* alive with cached/independent data far more easily than you can keep *checkout* alive — so protect them differently (Section 5).

2.4 The next day — Klaviyo (June 4, 2026)

Less than 24 hours later, **Klaviyo** — a marketing/CDP platform in the *event-ingestion and messaging* layer — had its own disruption. Per status reporting, the impact spanned **access to klaviyo.com, the APIs, event ingestion, and subscriptions**: customers could be **unable to log in, send campaigns or flows, or have their events processed**. (The exact incident window isn't fully detailed in public reporting; treat the duration as "multi-hour, plan for the worst.")

Why this matters for the architecture:

- **It's a different layer failing.** Shopify is the *commerce channel*; Klaviyo is a *destination for customer events*. Two different platforms, two different layers, two consecutive days. The lesson generalizes: **assume every external dependency will fail; design so none of them can take you down.**
- **Event ingestion failing is a silent data-loss risk.** If your app fires customer/order events straight at Klaviyo and Klaviyo is rejecting them, those events are **gone** unless you buffered them. This is exactly why cross-system events must be **queued, archived (R2), and replayable** (Section 4) rather than fired-and-forgotten — when Klaviyo recovers, you replay the backlog and lose nothing.
- **Marketing/consent events belong in your data layer first.** If an event matters (a signup, a consent change, a purchase), it should land in **Xano** as the system of record and be *projected* to Klaviyo/GA4/Adobe — never originate only inside a third-party platform that can drop it.

2.5 Sources

- [Shopify Resolves 2-Hour Outage Impacting Storefronts and Checkouts — PYMNTS](#)
 - [Ecommerce Outage Playbook: Survive Platform Downtime — Digital Applied](#)
 - [Klaviyo Status — Incident History](#)
-

3. The ORM / Content / Data Layer split


"ORM" here doesn't mean a code library — it means the **object-relational mapping between systems**: how a product, customer, or order in one platform

is represented and kept consistent in the others. Get this mapping wrong and every outage, schema change, or rate-limit becomes a data-integrity incident.

3.1 Three layers, three owners

LAYER	OWNER	LIVES IN	MAPPING RESPONSIBILITY
Data layer (system of record)	Xano	Relational tables: catalog, variants, customers, consent, orders, sync state	Canonical IDs. Every external object maps back to a Xano row by a stable natural key.
ORM / projection layer	Cloudflare Worker	Stateless transforms + edge cache (KV)	Translates a Xano record $\Rightarrow$ Shopify object $\Rightarrow$ Webflow CMS item. Owns idempotency and conflict resolution.
Content / presentation layer	Webflow (+ Shopify storefront)	CMS items, pages, product display	Renders what the projection layer publishes. Never the origin of truth.

3.2 Direction of flow (the rule that survives outages)



The invariant: data flows **out of** Xano to every channel. No channel is allowed to be the only place a fact exists. Shopify receives the catalog and returns

orders; it does not *own* the catalog. When Shopify disappears for two hours, the catalog still exists in Xano, the content still renders from Webflow + edge cache, and orders captured during the gap reconcile when Shopify returns.

3.3 Mapping discipline (avoid the classic ORM bugs)

- **Stable natural keys**, not platform IDs, as the join key (e.g. SKU/handle), so a record survives being re-created in any one platform.
- **Idempotent upserts** keyed on `(natural_key, source)` so a retried webhook or a replayed event can't create duplicates.
- **One writer per field**. Decide which platform is authoritative for each field (price → Xano; published-state → Webflow; fulfillment-state → Shopify) and never let two systems write the same field.
- **Versioned records** with `updated_at` / monotonic version so conflict resolution is deterministic (last-writer-wins by version, not by wall clock).

4. Automation recommendation — event-driven, not platform-coupled

4.1 Replace "live pull" with "cached read + async write"

The outage-resilient pattern, and the one this stack already implements:

PATTERN	READ PATH	WRITE PATH
× Fragile (don't)	Page calls Shopify live on every request	Page writes directly to Shopify, fails if Shopify is down
✓ Resilient (do)	Page reads from Xano / Cloudflare edge cache; Shopify is refreshed asynchronously	Writes go to a queue ; the Worker delivers to Shopify with retries, dead-letter, and circuit breaking

4.2 The delivery guarantees that make an outage survivable

These are operating defaults, not aspirations:

- **Queue every cross-system change** (`integration_queue` in Xano) instead of calling platforms inline. An outage just means the queue gets deeper, not that work is lost.
- **Per-priority retry with backoff** — critical inventory retries hardest, batch analytics softest:

PRIORITY	MAX RETRIES	BACKOFF	DEAD-LETTER AFTER
1 — inventory	5	10s · 30s · 2m · 10m · 30m	30 min
2 — orders	4	30s · 2m · 10m · 1h	1 hour
3 — customers	3	2m · 15m · 1h	1 hour
5 — analytics	2	1h · 6h	6 hours

- **Circuit breaker per destination** — if Shopify fails 5× in 10 minutes, stop hammering it for a cooldown, then auto-probe. During the June 3 outage this is what stops you from burning your rate limit and your logs against a dead endpoint, and lets you recover cleanly the instant it returns.
- **Dead-letter queue, never discard** — failed events are stored in Xano, archived in R2, and alerted on. After recovery you replay them in order. *Nothing captured during the outage is lost.*
- **A durable event archive (R2)** so you always have an independent, append-only record of what happened — including everything that piled up while a platform was down.

4.3 What stays on a timer

Cron is reduced to a **safety net**, not the primary mover: token refresh, dead-letter sweep, and a catch-up reconciliation pass that compares Xano ⇔ each

channel and heals drift introduced during an outage.

5. Outage runbook — what to do when a platform goes dark

This is the operational core of "process management." Keep it short, rehearsed, and owned.

5.1 Detection (minutes 0–5)

- **Watch external signal, not just your own dashboards.** A platform outage won't show as *your* error — it shows as upstream 5xx/timeouts and a DownDetector spike. Subscribe to each platform's status page.
- **Confirm scope:** is it read, write, or both? Storefront, admin, checkout, POS, or API? (June 3 hit *all* of them.)
- **Confirm blast radius:** which of your surfaces actually depend on the down platform right now?

5.2 Containment (minutes 5–20)

IF DOWN...	DO
Shopify storefront	Serve cached catalog/content from Cloudflare + Webflow. Replace any Shopify-rendered surface you control with your own honest outage banner — never let customers land on <i>"This store does not exist."</i>
Shopify checkout / payments	Show a clear "checkout temporarily unavailable — we'll email you / try again shortly" state. Capture intent (email, cart) into Xano so you can recover the sale. Do not silently fail the buy button.
Shopify admin / API	Stop inline writes; let the queue absorb them. Verify the circuit breaker has opened.
Xano (source of truth)	This is your most severe case — failover/read-replica strategy and edge cache become primary. Halt writes that can't be reconciled.
Webflow	Serve last-published static content from cache/CDN; pause CMS publishes.
Cloudflare	Lean on edge cache + status comms; this is the platform whose job is to <i>be</i> the resilient layer, so its own incidents are highest-severity.

5.3 Communicate (in parallel)

- Post your own status (status page / banner / social) **before** customers post for you. "We're aware, here's what works, here's what doesn't, here's the ETA."
- Brief support with the known facts and the workaround.

5.4 Recovery (when the platform returns)

1. **Don't trust "resolved."** Probe with a single canary request before reopening the floodgates.
2. **Let circuit breakers auto-close**; watch error rate as traffic resumes.
3. **Replay the dead-letter queue** in priority order.
4. **Run reconciliation** (Xano $\rightleftharpoons$ Shopify $\rightleftharpoons$ Webflow) to heal any drift — especially orders/inventory that moved during the gap.

5. **Remove your outage banners** only after reconciliation is clean.

5.5 Post-incident (within 48h)

Short, blameless write-up: timeline, customer impact, what the cache/queue saved you, what didn't degrade gracefully, and one or two concrete hardening actions. File it; review it next time.

6. Roles & responsibilities (RACI-lite)

ACTIVITY	PLATFORM ENG	OPS / SUPPORT	BUSINESS OWNER
Schema / mapping changes (Xano source of truth)	Owns	Informed	Consulted
Channel config (enable/disable Shopify, GA4, Adobe...)	Consulted	Owns	Informed
Deploys (Worker / extension / app)	Owns	Informed	Informed
Outage detection & runbook execution	Owns	Owns (comms)	Informed
Customer communication during incident	Consulted	Owns	Accountable
Post-incident review	Owns	Contributes	Accountable

Principle (from the security architecture): each platform is independently controllable. Disabling a misbehaving channel is a **config toggle, not a code change** — so Ops can contain an incident in seconds without waiting on a deploy.

7. Change & deploy process (summary)

The stack has **three independently deployable units** — they ship separately so a change to one can't take down the others:

1. **Cloudflare Worker** — the runtime/ORM-projection layer (data plane).
2. **Webflow Designer Extension** — the operator interface (validate before publish).
3. **Public spec / docs** — stakeholder-facing, **no source code, no secrets**.

Change-management rules:

- **One source of truth per fact** — changes to canonical data go through Xano; channels are projected, never hand-edited into drift.
- **Backward-compatible first** — add new endpoints/aliases before redirecting traffic; run old and new paths in parallel, then retire.
- **Every config change is authenticated and logged** with before/after, so any change can be audited and reversed.
- **Never commit secrets**; credentials live only in the encrypted per-tenant store, never in code or docs.

8. Checklist — is your stack June-3-proof?

- ☐ **Xano is the source of truth** — every catalog/customer fact exists independently of Shopify.
- ☐ **Storefront reads survive a Shopify outage** from Xano + edge cache.
- ☐ **You control the outage page** — customers never see *"This store does not exist."*

- ❑ **Checkout fails loudly and captures intent** instead of silently dropping sales.
 - ❑ **All cross-system writes go through a queue** with retry + dead-letter + circuit breaker.
 - ❑ **Marketing/event sends (Klaviyo, GA4, Adobe) are buffered**, not fired-and-forgotten — a downstream outage backs up the queue instead of losing events.
 - ❑ **Nothing is discarded on failure** — dead letters are stored, archived, and replayable.
 - ❑ **A reconciliation pass** can heal Xano $\rightleftharpoons$ Shopify $\rightleftharpoons$ Webflow drift after recovery.
 - ❑ **Status-page subscriptions** for all four platforms feed your alerting.
 - ❑ **The outage runbook is written, owned, and rehearsed** — not improvised at 9:27 a.m.
 - ❑ **One writer per field**; mapping uses stable natural keys and idempotent upserts.
-

9. Verification — the process has to leave evidence

Everything above keeps the stack *running*. This section is about being able to **prove afterwards what it did** — which is now a separate requirement, and one an outage runbook alone does not satisfy.

The distinction that matters: a log records that something happened; **evidence lets a third party confirm it without trusting you**. Only the second survives an auditor, a regulator, or a counterparty in dispute.

What this stack publishes, all public and unauthenticated:

SURFACE	PURPOSE
<code>/.well-known/jwks.json</code>	The Ed25519 public key set. The anchor for every signature below.
<code>/license/verify</code>	Check any certificate this platform issued — no account, no call to us.
<code>/v/<record-hash></code>	Short verification link; resolves the full signed certificate.
<code>/license/qr/<hash>.svg</code>	Print-ready QR to the verification page — for packaging, labels, documents.

Process rules that make the evidence hold:

- **Append-only.** Corrections add a row; they never edit one. A history an administrator can silently amend proves nothing, including the parts that are correct.
- **Record denials too.** A refused action is evidence the controls were operating. A log containing only successes cannot demonstrate a control at all.
- **One writer per field, idempotent upserts** — already the rule in §3.3. It is also what makes the ledger reconstructible after an outage rather than merely plausible.
- **Rotate keys forward.** Mint a new pair, repoint the active identifier; certificates issued earlier stay verifiable against the published retired key. A suspected compromise becomes a pointer change, not an identity loss.

Deeper treatment: [Agent Authority — Technical Brief](#) · [AI Trust Framework Requirements](#) · [The Trust Framework \(non-technical\)](#).

10. SBOM and firmware — the same discipline, applied to what you ship

If the stack ships anything with digital elements, the EU Cyber Resilience Act turns this from good practice into market access.

The dates that drive planning: CRA entered into force 10 December 2024.

Article 14 reporting obligations apply from 11 September 2026 — actively exploited vulnerabilities and severe incidents, reported to ENISA and national CSIRTs, *including for products already on the market*. Full obligations — secure update mechanisms, conformity assessment, CE marking — apply 11 December 2027.

Read the September 2026 date carefully, because it is the one that changes engineering priorities: **you cannot report exploitation you cannot see**. A reporting deadline is, in practice, an evidence-ledger deadline.

Process shape:

- **Firmware is a grant, not a URL.** A signed installer behind a storage link is, once shared, shared forever, with no record of who fetched it or when. Vault the image; release it through an authorization gate.
- **Three records outlive every transaction** — a registry row (product, version, size, full SHA-256, linked SBOM), an upload certificate (signed proof of what was vaulted, by whom, when), and a hash-chained access ledger where every grant, every denial, and every served byte-stream seals the row before it.
- **The SBOM is the dependency inventory** you will be asked for during an incident. Generate it at build time and store it with the artifact — reconstructing it later is exactly the archaeology §9 exists to prevent.

Detail: Firmware, SBOM & the Cyber Resilience Act.

11. AEO — make the process machine-readable

The same content that documents your stack for people now has a second audience: retrieval systems, agents and answer engines. This is not marketing.

It determines whether an operator asking a question at 2 a.m. — or an agent acting on their behalf — finds the correct answer or an inference.

Surfaces this stack publishes:

- `/docs-index.json` — the machine index of every published document, with canonical render URLs.
- `llms.txt` — the corpus map for language models.
- **JSON-LD on every doc and product** — `DefinedTermSet` for terminology, structured product data carrying the same GTIN/MPN the catalog publishes.
- **RAG-backed search** at `/pages/knowledge-base` — answers cite their sources, so a claim can be traced back to a document rather than trusted.

Process rules:

- **One canonical page per topic.** Duplicates split retrieval and produce confidently wrong answers.
- **Publish to every surface or none.** A document in the index but missing from the CMS produces a citation that deep-links to nothing — a silent failure, and one this project has shipped before.
- **Identifiers must agree across planes.** The MPN in the ledger, the MPN in the product page's structured data, and the MPN in the feed are the same string or the join is fiction.
- **Terminology is infrastructure.** Define the terms your team will be tested on — idempotent and monotonic are the two a QA function cannot operate without.

12. Extended checklist — evidence, artifacts, retrieval

- ☐ **A third party can verify a record** without your cooperation and without an account.
- ☐ **Records are append-only**; corrections add rather than edit.
- ☐ **Denials are recorded**, not just successes.
- ☐ **Keys are rotatable** without invalidating certificates already issued.
- ☐ **SBOM is generated at build time** and stored with the artifact.
- ☐ **Firmware releases through an authorization gate**, never a durable link.
- ☐ **An incident report can be produced from the ledger**, not reconstructed from exports.
- ☐ **Every published doc exists on all surfaces** — repo, index, CMS, KB page, RAG corpus.
- ☐ **Product identifiers agree** across ledger, structured data, and channel feed.
- ☐ **Retries are idempotent and sequences monotonic** — duplicates blocked, ordering provable.

*If you can check every box, a global Shopify outage is a degraded hour, not an existential one. That is the entire point of running four platforms instead of one. And if you can check the boxes in §12, the hour is also **documented** — which is the difference between an incident you explain and one you merely survive.*
