

# Google turned off the list upload. We were never uploading lists.

On April 1, 2026, Google disabled Customer Match uploads through the Ads API and replaced them with the Data Manager API — an ingest channel that **requires** per-member consent fields on the wire. For a pipeline that always computed segments server-side and gated them on consent, nothing had to be unlearned. This plan connects Shopify segments to Google along three planes: **Reach** (value-weighted audiences), **Memory** (a BigQuery history), and **Answers** (plain-English questions over both).

For: Marketing ops, analytics engineering, and the business analyst who owns segments

Status: Plan of record · GA4 pipe built · Live Google push routes via the Data Manager API

Date: 2026-07-10   Depends on: [SEGMENTS-GA4-BIDDING.md](#) · [ARCHITECTURE.md](#) · [GLOBAL-PAYOUTS.md](#)

---

## 0. TL;DR

---

On **April 1, 2026**, Google disabled Customer Match uploads through the Google Ads API for developer tokens without prior Customer Match history. The replacement — the **Data Manager API** — needs no developer token, authenticates with a plain Google Cloud service account, and fans one ingest endpoint out to Google Ads, Display & Video 360, and GA4. For a pipeline that was always consent-gated and server-computed, this is not a migration. It is the arrival of the channel the architecture assumed.

This plan connects Shopify segments to Google along **three planes**: **Reach** (segments become value-weighted ad audiences), **Memory** (segment, consent, and revenue history lands in BigQuery), and **Answers** (Google's Conversational Analytics sits on that warehouse so a non-technical stakeholder can ask "which segments drove revenue this week?" in plain English).

**One-line recommendation:** *Ship the single-US-market pipe now on the Data Manager API, with a `market` key on every row from day one — so going global is adding markets, not rebuilding.*

---

1. WHAT CHANGED AT GOOGLE – AND WHY IT FAVORS THIS DESIGN

| RETIRED (APRIL 1, 2026)                                                                      | REPLACEMENT                                                                                                        |
|----------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------|
| Google Ads API <code>OfflineUserDataJobService</code> / <code>UserDataService</code> uploads | <b>Data Manager API</b> – <code>datamanager.googleapis.com</code>                                                  |
| Developer token + Ads-specific OAuth plumbing                                                | Standard Google Cloud OAuth; one scope ( <code>auth/datamanager</code> )                                           |
| Per-product upload endpoints                                                                 | One ingest endpoint → Google Ads, DV360, GA4/Firebase, Ad Manager, CM360, SA360                                    |
| Consent as a policy footnote                                                                 | Consent fields ( <code>ad_user_data</code> , <code>ad_personalization</code> ) <b>on every member, on the wire</b> |

The Data Manager API also supersedes the GA4 Measurement Protocol for server-side conversion ingestion. The GA4 pipe described in `SEGMENTS-GA4-BIDDING.md` keeps working – this plan adds the direct channel beside it.

The consent posture doesn't change at all: the pipeline computes segments in the data layer, gates them on consent, and hands Google a qualified signal – never a raw list. Google's new API now *requires* what this architecture already *did*.


2. WHAT COUNTS AS A SEGMENT

Four kinds of group, and they travel differently. Only one ever becomes an advertising audience.

| DOMAIN             | EXAMPLE                                         | WHERE IT GOES                                                                                                                   |
|--------------------|-------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------|
| Peers & Households | VIP repeat buyers; a family sharing one account | The <b>only</b> domain projected to Google Ads – consented members, hashed identifiers, value-weighted                          |
| Teams              | A design team; an approvals group               | Insight plane only (warehouse + answers). Never advertising                                                                     |
| Organizations      | A client company; an agency workspace           | Insight plane; any B2B projection is a separate, later decision                                                                 |
| Nations            | Canada; the EU; Korea                           | Not an audience – a <b>boundary</b> . Scopes campaigns, partitions the warehouse, selects the consent regime and data residency |

**Rule of thumb:** Peers/Households flow *out* to Google; Teams and Organizations flow *up* to insight; Nations decide *where* – and under *which law* – everything runs.

### 3. THE THREE PLANES



| PLANE   | WHAT IT DELIVERS                                                                                     | WHO IT SERVES            |
|---------|------------------------------------------------------------------------------------------------------|--------------------------|
| Reach   | Consent-gated, revenue-weighted Customer Match audiences + value-based conversions for Smart Bidding | Performance media        |
| Memory  | Segment / consent / revenue history in BigQuery – atomic nightly loads, every run in the audit chain | Analytics engineering    |
| Answers | Natural-language questions over the warehouse – no SQL, no analyst queue                             | The BA who owns segments |

The Answers plane is the accessibility thesis made literal: connected CRM data that a non-technical teammate can interrogate directly.

#### 4. GOOGLE-SIDE REQUIREMENTS

One Google Cloud project. Services to enable:

| API                                                                                            | WHY                                                                           |
|------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------|
| <code>datamanager.googleapis.com</code>                                                        | Audience ingestion (Customer Match successor) + weighted conversion ingestion |
| <code>analyticsadmin.googleapis.com</code>                                                     | Programmatic GA4 audiences mirroring segment definitions                      |
| <code>bigquery.googleapis.com</code> + <code>storage.googleapis.com</code>                     | The warehouse and its staging bucket                                          |
| <code>geminidataanalytics.googleapis.com</code> + <code>cloudaicompanion.googleapis.com</code> | Conversational Analytics data agents (Answers plane)                          |

**Authentication:** a service account, not user OAuth. Service accounts skip Google's OAuth app-verification queue; the runtime mints short-lived access tokens server-side (JWT-bearer exchange, cached until a minute before expiry). The credential enters the system through the same supervised key ceremony as every other secret — see `KEY-MANAGEMENT-LIFECYCLE` .

**Account linking (human, one sitting):** grant the service-account identity access to the destination Google Ads account (requires Ads admin), link GA4 ↔ Ads, and verify **Customer Match eligibility** — it is account-level (policy compliance + payment history) and worth confirming *before* anything is built on top.

**Consent on the wire:** every ingest carries per-member `ad_user_data` and `ad_personalization` , mapped from the same consent records that already gate the GA4 pipe. No consent, no export — structurally, not procedurally.

Setup is a CLI-and-console sitting; the runtime adds **no new infrastructure dependency** — the existing worker and data layer call Google's REST endpoints directly.

#### 5. THE WAREHOUSE FEED

Deltas come off a **sync queue**, not a timestamp watermark — the queue gives idempotency, retry, and dead-lettering, and keeps the feed inside the audit chain (watermarks silently miss

hard-deletes).

| INGEST OPTION                        | FIT                                                                                                                                           |
|--------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------|
| GCS batch load + MERGE – recommended | Ingestion is free and atomic (a night's load lands completely or not at all); true upserts against the system of record                       |
| Legacy streaming insertAll           | Sub-minute freshness, but billed at a 1 KB-per-row minimum and rows linger in a streaming buffer                                              |
| Storage Write API                    | The best engine, but gRPC – reachable from an HTTP-only backend via a thin edge shim, only worth it if a stakeholder needs sub-minute numbers |

The loop: read pending → batch to NDJSON → load into staging → MERGE on the business key → advance the queue **only on success**. Failures dead-letter with the job ID; every run logs counts into the audit chain.

## 6. ONE MARKET OR WORLDWIDE

Google's fees are near zero either way at this volume – batch loads are free, and the free tiers (10 GiB storage, 1 TiB query per month) cover the working set. The real cost axis is accounts, compliance, and people-time.

|                | SINGLE US MARKET                                       | GLOBAL ORGANIZATION                                                                                             |
|----------------|--------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------|
| Ads structure  | One Ads account                                        | Manager account (MCC) + per-market accounts                                                                     |
| GA4            | One property                                           | Per-market properties, each linked to Ads                                                                       |
| Warehouse      | One US-region dataset                                  | Region-partitioned datasets (EU data stays in the EU)                                                           |
| Consent regime | US opt-out rules; Consent Mode v2 already exceeds them | EEA: all four Consent Mode v2 signals <b>mandatory</b> (built); Korea and Japan add their own consent paperwork |
| Nations domain | Dormant                                                | Active — where country boundaries do real work                                                                  |
| Google cost    | ≈ \$0 / month                                          | ≈ \$0–10 / month                                                                                                |
| Real cost      | One setup sitting; live in days                        | Per-market account grants + legal review — people-time, not fees                                                |

**Recommendation: A with a B-shaped schema.** Ship the single US market now; carry a `market` key on every warehouse table and campaign row from day one, and name datasets so regional siblings can appear beside them. Going global then reuses the market plumbing the Canada proof-of-concept already validated.

7. DELIVERY PHASES

| PHASE         | DELIVERS                                                                 | VISIBLE RESULT                                    |
|---------------|--------------------------------------------------------------------------|---------------------------------------------------|
| 0 – Setup     | Cloud project, APIs, service account, Ads + GA4 links, eligibility check | A verified checklist; nothing live yet            |
| 1 – Feed      | Nightly queued delta load into BigQuery                                  | Row counts + "last updated" any morning           |
| 2 – Reach     | Live audience ingest via Data Manager API (test mode stays the default)  | The segment appears in Google Ads with a count    |
| 3 – Freshness | Event-triggered incremental updates (no polling)                         | Audience counts move on their own after busy days |
| 4 – Mirror    | GA4 audiences created programmatically per segment definition            | Event-based audiences beside Customer Match       |
| 5 – Value     | Weighted conversions ingested – the Smart Bidding payoff                 | Conversion value reflects segment weights         |
| 6 – Answers   | Conversational Analytics over the warehouse                              | A plain-English question box for segments         |

Phase 0 is a supervised human sitting. Phases 1 and 2 are independent and parallel. Phase 6 waits only on Phase 1.

8. THE ENTERPRISE PIVOT – WHEN THE INCUMBENT STACK CAN'T CONFORM

Enterprise organizations holding seven-figure Salesforce/Adobe commitments are waiting for those platforms to conform to the new advertising regime. They won't – not because the vendors are slow, but because the regime broke the architecture those stacks were built on:

- **Consent Mode v2** demands per-event, per-signal consent stamped at the data layer. Incumbent CDPs store consent as a field on a record – the wrong place architecturally, and no release cycle relocates it.
- **The Customer Match retirement** (April 1, 2026) deleted the upload pattern every enterprise Google connector was built around. The successor channel expects consent

fields on every member, on the wire — which presumes the gate above already exists.

- **Agentic buyers never fire the instrumentation.** An agent completes search → cart → checkout with no click, no page view, no pixel. A seven-figure measurement estate is not underperforming on this channel — it is structurally blind to it.

## The CDP that syndicated the page view

The clearest instance is the Segment-class CDP. Its founding primitives are the Universal Analytics worldview promoted to infrastructure: `page()` as a first-class API call, identity rooted in a cookie-scoped `anonymousId`, collection defaulting to a client-side library injected into the DOM. Instead of sending the page view to one destination, it sold sending it to a hundred — so it did not escape the page view's obsolescence, it **syndicated** it: when the cookie/client-script architecture broke, the value proposition broke in every destination simultaneously. Server-side sources and audience products moved the transport, not the worldview — the same page/track/identify schema, the same cookie-rooted identity spine, audiences computed inside the toll booth and synced to ad platforms through exactly the upload patterns retired in April 2026. The pricing makes the toll explicit: metering by monthly tracked users bills the merchant per anonymous cookie, including the overwhelming share who never consent, never sign in, and never buy.

This substrate inverts every term of that model. Consent gates the signal at the source instead of filtering it at the destination; identity is the consented login, not a stitched cookie; audiences are computed in the merchant's own data layer and projected outward; and nobody is billed per ghost.

The pivot is not rip-and-replace. The incumbent stack stays as the system of engagement, and keeps receiving its feeds. What changes is the signal path: identity, consent, mandates, and conversions run through the consent-gated substrate this plan describes — stamped at the source, encrypted in transit, agent-addressable — and the estate consumes from a plane that is admissible under the new rules.

**Your stack isn't wrong — it's deaf to the new signals. We don't replace it; we give it ears that are legal to use — and ears an average user in the organization can operate.** The substrate ships with its own plain-language surfaces: a question box that answers in English (the Answers plane above), a searchable documentation archive written for operators — this article is part of it — and a campaign wizard with approval gates. The people who own segments day-to-day run this in-house. Conformance here is not a statement of work, and nothing waits on a consulting team ten time zones away.

---



## 9. RELATED DOCUMENTS

---

- `SEGMENTS-GA4-BIDDING.md` — the built GA4 pipe: consent gate, revenue weights, user properties, audiences. This plan's Reach plane is its direct-channel sibling.
- `ARCHITECTURE.md` — where the worker, data layer, and consent plane sit.
- `KEY-MANAGEMENT-LIFECYCLE` — the ceremony that admits the service-account credential.
- `GLOBAL-PAYOUTS.md` — the settlement side of the same market boundaries the Nations domain scopes.